

Intel DPDK

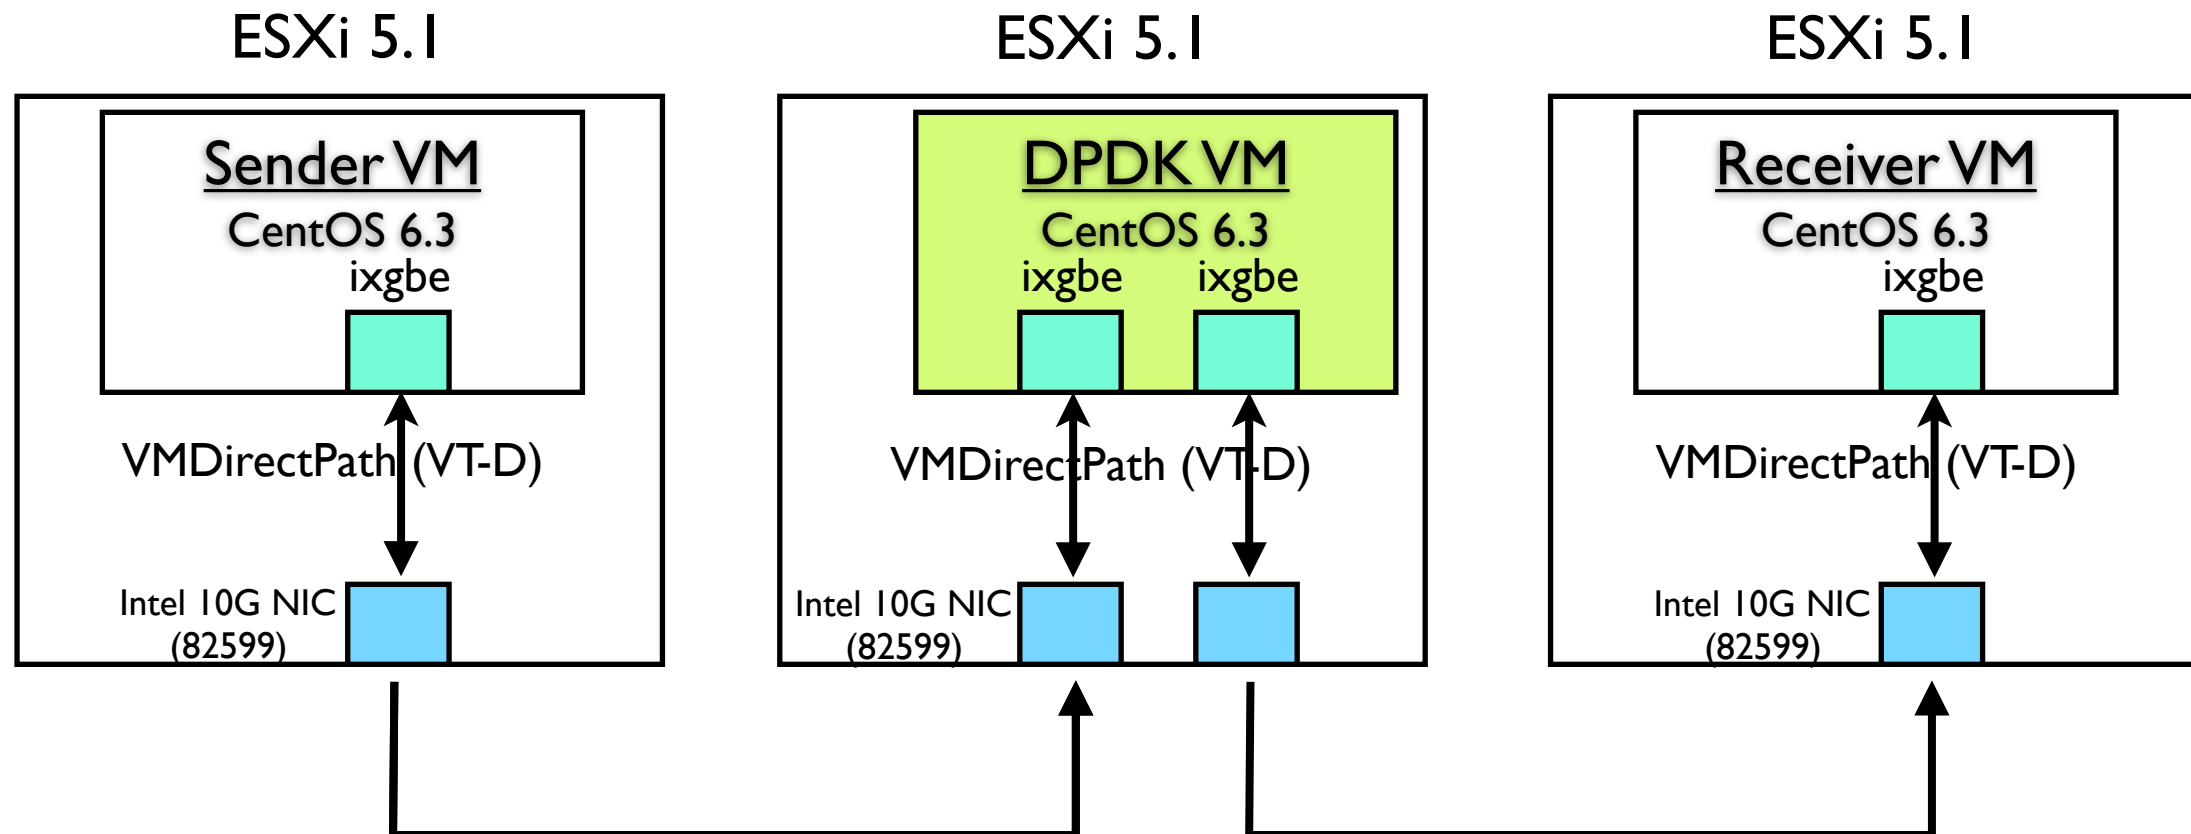
Step by Step Instructions

Hisaki Ohara (@hisak)

Objectives

- Build/Execute sample applications (helloworld, L2fwd and L3fwd)
- Packet forwarding by generating with Linux/pktgen

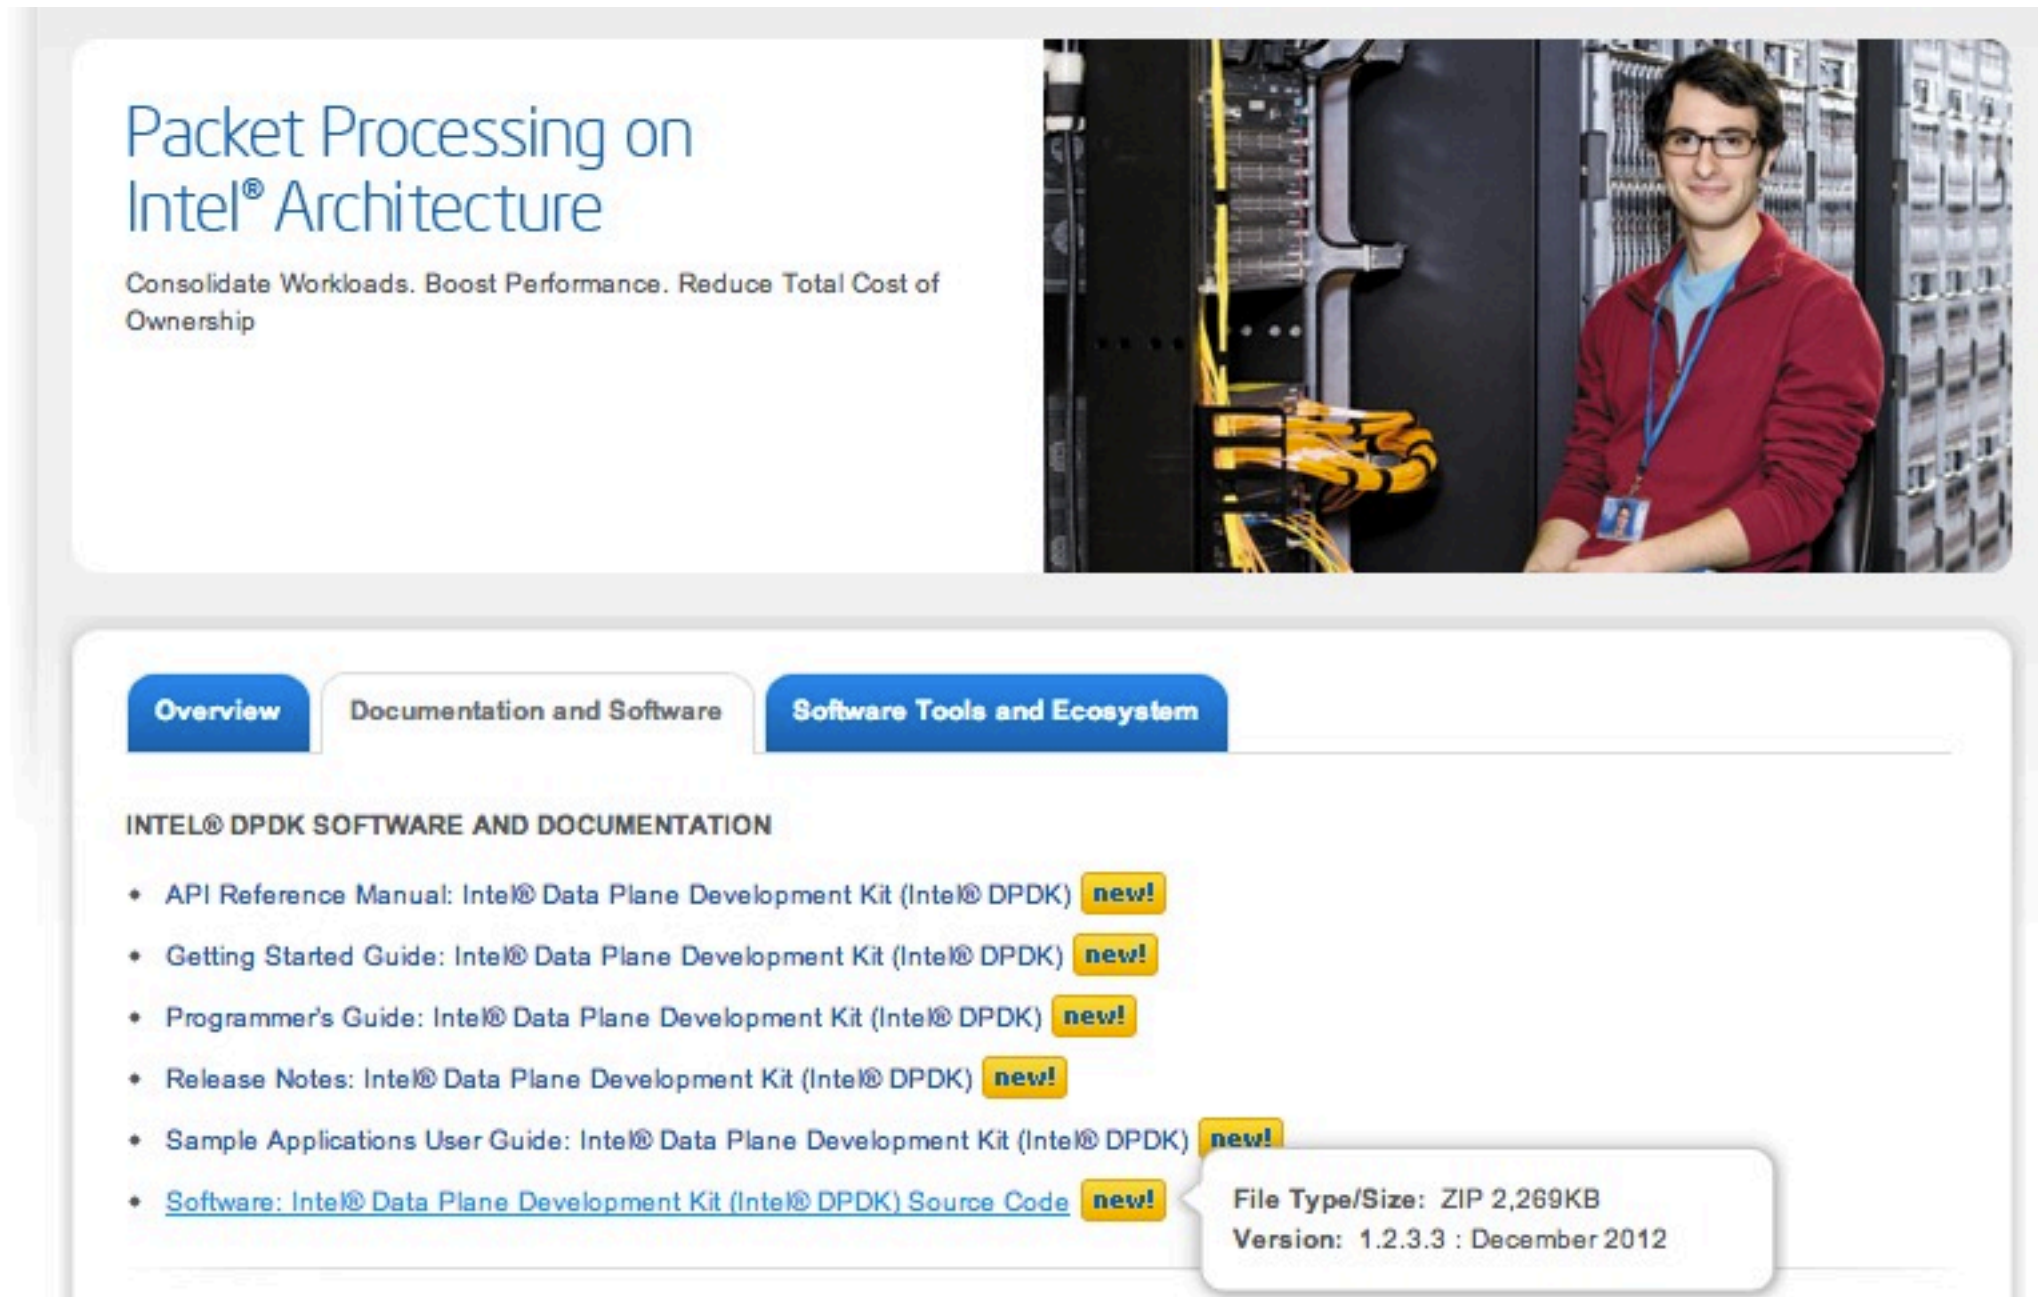
Test Environment



- ESXi 5.1
 - CPU: Xeon 5600 Series
 - Guest OS: CentOS 6.3 x86_64
 - # of vCPUs: 2
 - 10G NIC (82599) is passed through
 - In-box driver of ixgbe

Step0: Download source codes

- Source codes and relevant documents
- <http://www.intel.com/go/dpdk>



The screenshot displays the Intel DPDK website. At the top, the heading 'Packet Processing on Intel® Architecture' is followed by the tagline 'Consolidate Workloads. Boost Performance. Reduce Total Cost of Ownership'. A photo of a man in a red jacket stands next to server racks. Below this, three tabs are visible: 'Overview', 'Documentation and Software', and 'Software Tools and Ecosystem', with the last one being active. Under the 'INTEL® DPDK SOFTWARE AND DOCUMENTATION' section, a list of resources is provided, each marked with a 'new!' badge. A tooltip for the 'Software: Intel® Data Plane Development Kit (Intel® DPDK) Source Code' item shows its file type and size as 'ZIP 2,269KB' and its version as '1.2.3.3 : December 2012'.

Packet Processing on Intel® Architecture

Consolidate Workloads. Boost Performance. Reduce Total Cost of Ownership

Overview Documentation and Software **Software Tools and Ecosystem**

INTEL® DPDK SOFTWARE AND DOCUMENTATION

- API Reference Manual: Intel® Data Plane Development Kit (Intel® DPDK) **new!**
- Getting Started Guide: Intel® Data Plane Development Kit (Intel® DPDK) **new!**
- Programmer's Guide: Intel® Data Plane Development Kit (Intel® DPDK) **new!**
- Release Notes: Intel® Data Plane Development Kit (Intel® DPDK) **new!**
- Sample Applications User Guide: Intel® Data Plane Development Kit (Intel® DPDK) **new!**
- [Software: Intel® Data Plane Development Kit \(Intel® DPDK\) Source Code](#) **new!**

File Type/Size: ZIP 2,269KB
Version: 1.2.3.3 : December 2012

Step I: Prepare Linux Kernel

- Add boot option and fstab for hugepage

```
# uname -a
Linux cent-dpdk 2.6.32-279.14.1.el6.x86_64 #1 SMP Tue Nov 6 23:43:09 UTC 2012 x86_64
x86_64 x86_64 GNU/Linux
# cat /boot/grub/grub.conf
<snip>
title CentOS (2.6.32-279.14.1.el6.x86_64)
    root (hd0,0)
    kernel /vmlinuz-2.6.32-279.14.1.el6.x86_64 ro root=/dev/mapper/vg_cent6-
lv_root rd_LVM_LV=vg_cent6/lv_swap rd_NO_LUKS LANG=en_US.UTF-8 rd_NO_MD
rd_LVM_LV=vg_cent6/lv_root SYSFONT=latarcyrheb-sun16 KEYBOARDTYPE=pc KEYTABLE=us
rd_NO_DM rhgb quiet crashkernel=auto hugepages=256
    initrd /initramfs-2.6.32-279.14.1.el6.x86_64.img
<snip>
# mkdir /hugepages
# cat /etc/fstab
<snip>
hugetlbfs                /hugepages                hugetlbfs rw,mode=0777 0 0
# reboot
```

- Confirm hugepage is enabled

```
# cat /proc/meminfo
<snip>
HugePages_Total:        256
HugePages_Free:         256
HugePages_Rsvd:         0
HugePages_Surp:         0
Hugepagesize:           2048 kB
<snip>
```

Step2: Build DPDK and samples

```
$ unzip INTELDPDK.L.1.2.3_3.zip
$ cd DPDK
$ make install T=x86_64-default-linuxapp-gcc
$ pwd
/home/dpdktest/DPDK
$ cd examples/helloworld
$ RTE_SDK=/home/dpdktest/DPDK make
  CC main.o
  LD helloworld
  INSTALL-APP helloworld
  INSTALL-MAP helloworld.map
```

Step3: helloworld sample

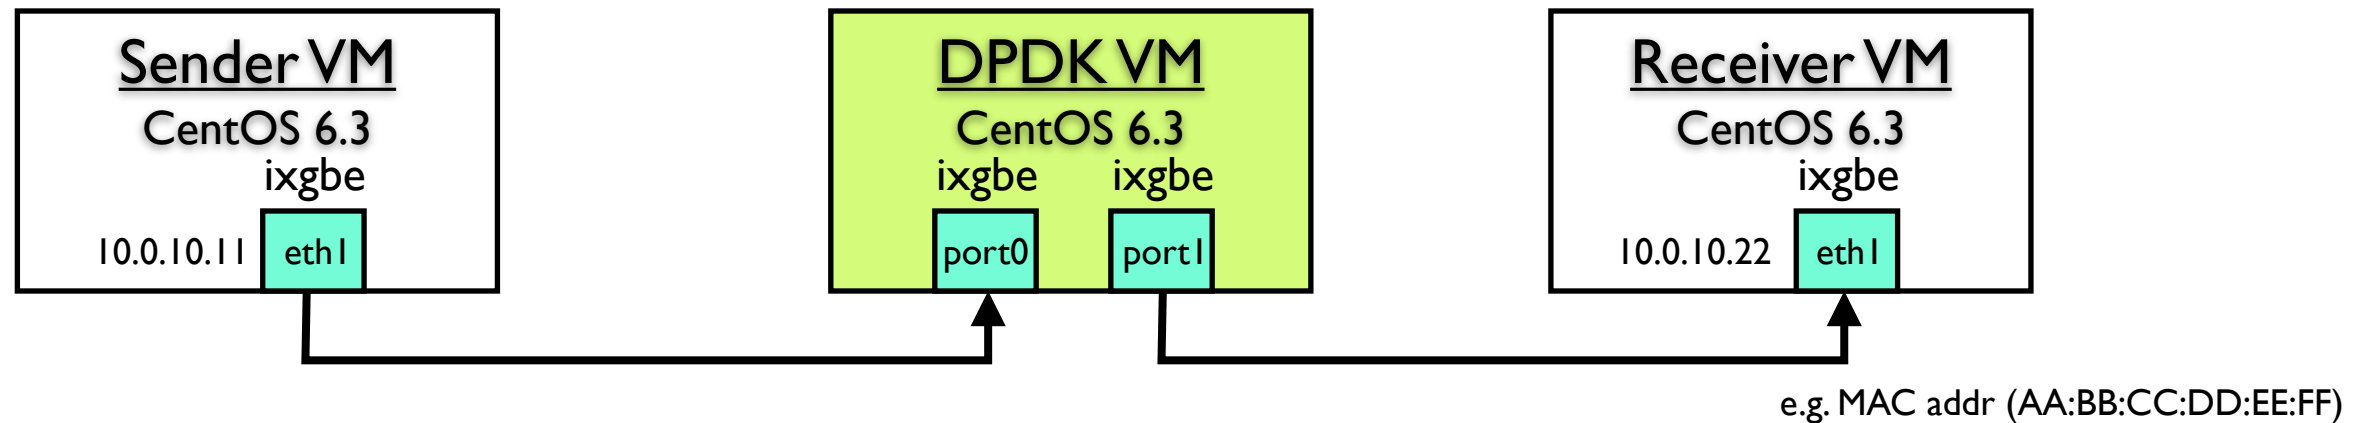
- Load required module for DPDK Linux app

```
# modprobe uio  
# insmod /home/dpdktest/DPDK/x86_64-default-linuxapp-gcc/kmod/igb_uio.ko
```

- Execute helloworld sample

```
# ./build/helloworld -c 3 -n 2  
EAL: coremask set to 3  
EAL: Detected lcore 0 on socket 0  
EAL: Detected lcore 1 on socket 0  
EAL: Requesting 256 pages of size 2097152  
EAL: Ask a virtual area of 0x20000000 bytes  
EAL: Virtual area found at 0x7f4862c00000 (size = 0x20000000)  
EAL: WARNING: Cannot mmap /dev/hpet! The TSC will be used instead.  
EAL: Master core 0 is ready (tid=82dd0800)  
EAL: Core 1 is ready (tid=621fe700)  
hello from core 1  
hello from core 0
```

Step4-1: L2fwd Sample Build

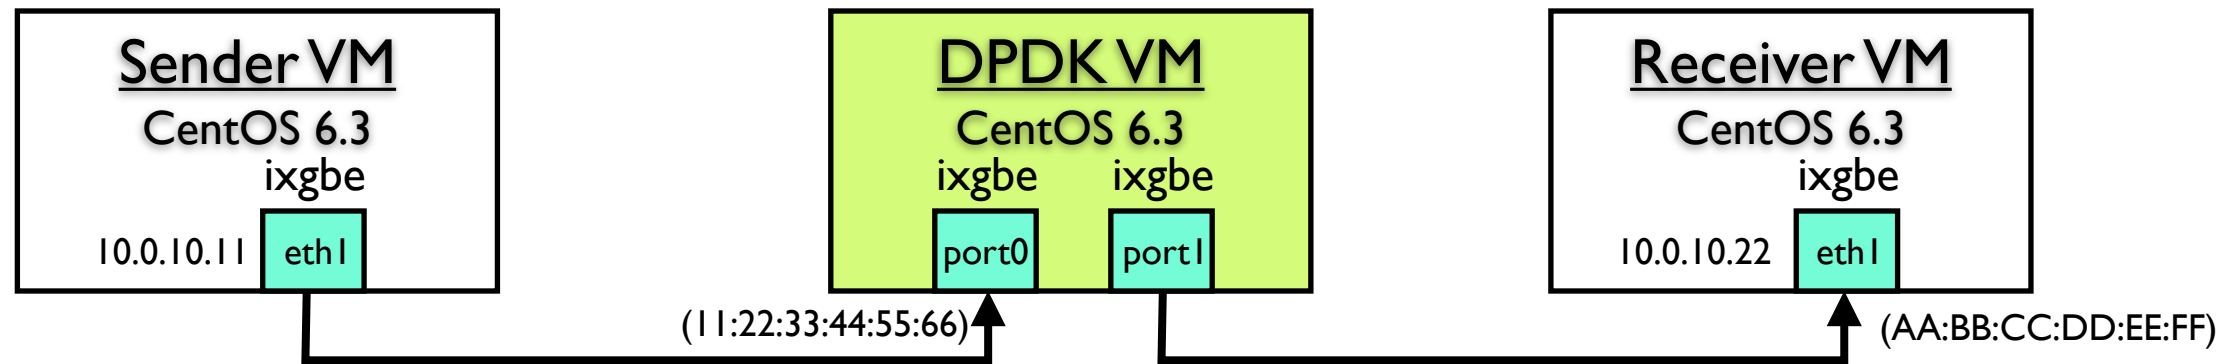


- L2fwd/L3fwd samples are very simple
 - One-way only. Don't expect ping/pong
 - Dest MAC address is hard-coded...

```
$ cd examples/l2fwd
$ diff -up main.c.0 main.c
@@ -293,7 +293,7 @@ l2fwd_simple_forward(struct rte_mbuf *m,
    /* 00:09:c0:00:00:xx */
    tmp = &eth->d_addr.addr_bytes[0];
-   *((uint64_t *)tmp) = 0x000000c00900 + (dst_port << 24);
+   *((uint64_t *)tmp) = 0xFFEEDDCCBBAA; /* AA:BB:CC:DD:EE:FF */

    /* src addr */
    ether_addr_copy(&l2fwd_ports_eth_addr[dst_port], &eth->s_addr);
$ RTE_SDK=/home/dpdktest/DPDK make
```


Step4-2: L2fwd Sample



```
#ip address add 10.0.10.11/24 dev eth1
#modprobe pktgen
```

```
#echo "rem_device_all" > /proc/net/pktgen/kpktgend_0
#echo "add_device eth1" > /proc/net/pktgen/kpktgend_0
#echo "count 10000000" > /proc/net/pktgen/eth1
#echo "clone_skb 1000000" > /proc/net/pktgen/eth1
#echo "pkt_size 60" > /proc/net/pktgen/eth1
#echo "delay 0" > /proc/net/pktgen/eth1
#echo "dst 10.0.10.22" > /proc/net/pktgen/eth1
#echo "dst_mac 11:22:33:44:55:66" > /proc/net/pktgen/eth1
#echo "start" > /proc/net/pktgen/pgctrl
```

```
#./build/l2fwd -c 0x3 -n 2 -- -p 0x3
```

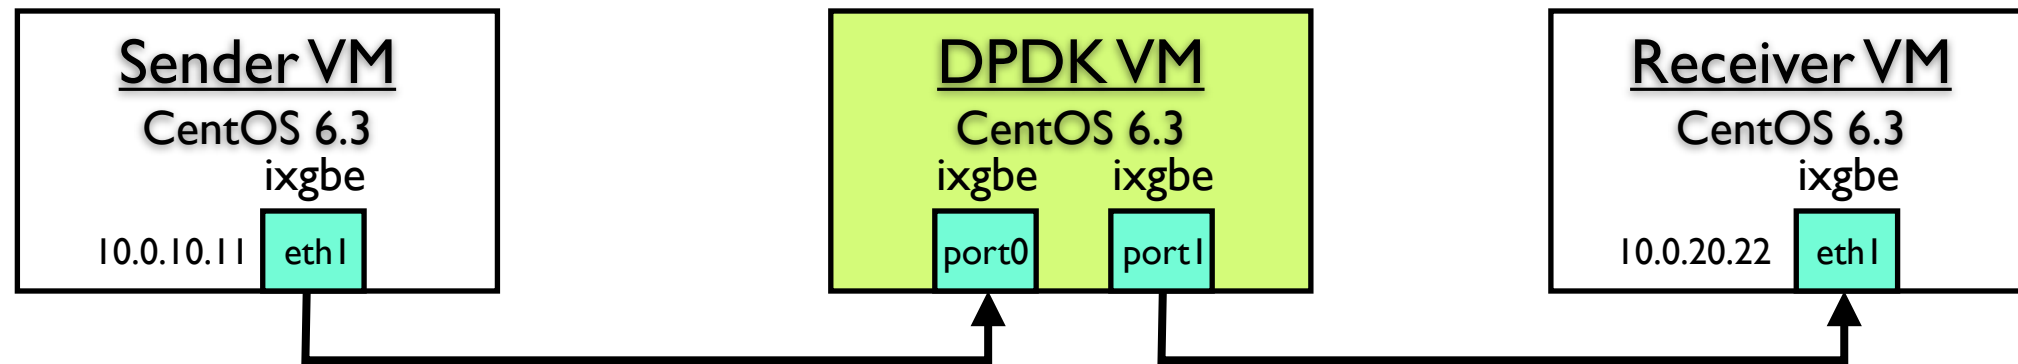
```
#ip address add 10.0.10.22/24 dev eth1
#vnstat -i eth1 -l
```

	rx		tx
bytes	150.70 MiB		0 KiB
max	370.17 Mbit/s		0 kbit/s
average	34.29 Mbit/s		0.00 kbit/s
min	0 kbit/s		0 kbit/s
packets	2633611		0
max	789696 p/s		0 p/s
average	73155 p/s		0 p/s
min	0 p/s		0 p/s

```
Port statistics =====
Statistics for port 0 -----
Packets sent: 0
Packets received: 5017107
Packets dropped: 0
Statistics for port 1 -----
Packets sent: 5017083
Packets received: 0
Packets dropped: 24
Aggregate statistics =====
Total packets sent: 5017083
Total packets received: 5017107
Total packets dropped: 24
=====
```

Packets are dropped at RX ports of DPDK VM and Receiver VM

Step5-1: L3fwd Sample Build



- L3 fwd sample has two functions to determine destination port
 - [default] destination IP address (LPM-based)
 - 5-tuples (Hash-based)

e.g. MAC addr (AA:BB:CC:DD:EE:FF)

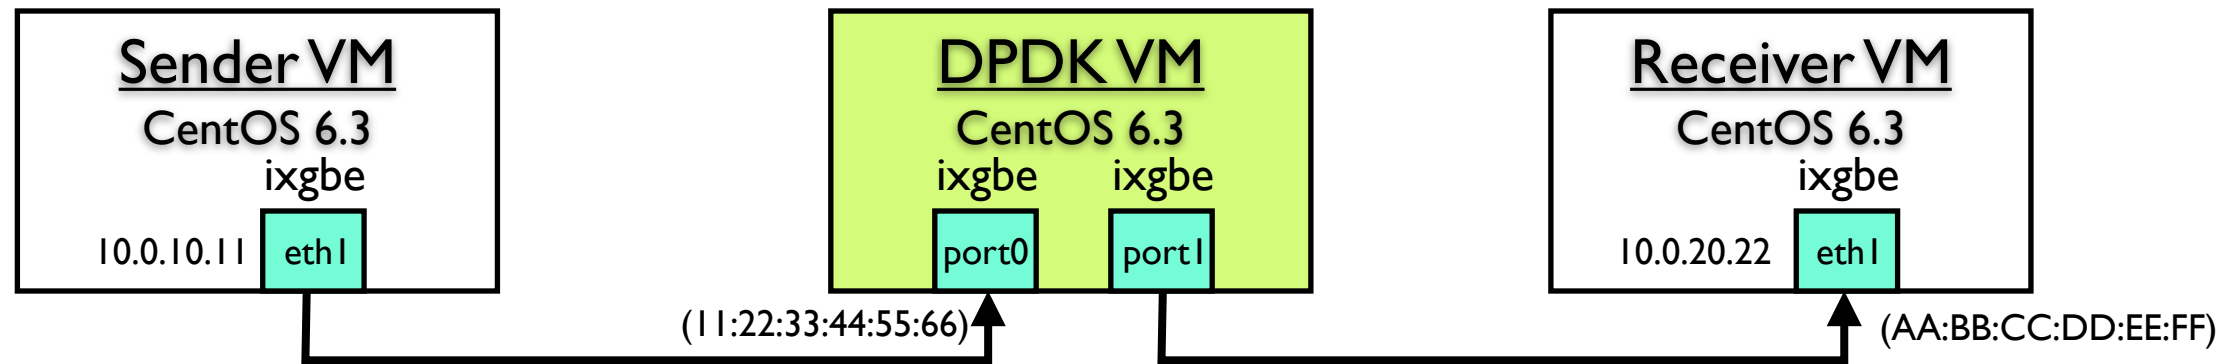
```
$ cd examples/l3fwd
$ diff -up main.c.0 main.c
@@ -282,6 +282,8 @@ static struct l3fwd_route l3fwd_route_ar
    {IPv4(6,1,1,0), 24, 5},
    {IPv4(7,1,1,0), 24, 6},
    {IPv4(8,1,1,0), 24, 7},
+   {IPv4(10,0,10,11), 24, 0},
+   {IPv4(10,0,20,22), 24, 1},
};

#define L3FWD_NUM_ROUTES \
@@ -475,7 +477,7 @@ l3fwd_simple_forward(struct rte_mbuf *m,

    /* 00:09:c0:00:00:xx */
    tmp = &eth_hdr->d_addr.addr_bytes[0];
-   *((uint64_t *)tmp) = 0x000000c00900 + (dst_port << 24);
+   *((uint64_t *)tmp) = 0xFFEEDDCCBBAA; /* AA:BB:CC:DD:EE:FF */

$ RTE_SDK=/home/dpdktest/DPDK make
```

Step5-2: L3fwd Sample



```
#ip address add 10.0.10.11/24 dev eth1
#modprobe pktgen
```

```
#./build/l3fwd -c 0x3 -n 2 -- -p 0x3
--config="(0,0,0),(1,0,1)"
```

```
#ip address add 10.0.20.22/24 dev eth1
#vnstat -i eth1 -l
```

```
#echo "rem_device_all" > /proc/net/pktgen/kpktgend_0
#echo "add_device eth1" > /proc/net/pktgen/kpktgend_0
#echo "count 10000000" > /proc/net/pktgen/eth1
#echo "clone_skb 1000000" > /proc/net/pktgen/eth1
#echo "pkt_size 60" > /proc/net/pktgen/eth1
#echo "delay 0" > /proc/net/pktgen/eth1
#echo "dst 10.0.20.22" > /proc/net/pktgen/eth1
#echo "dst_mac 11:22:33:44:55:66" > /proc/net/pktgen/eth1
#echo "start" > /proc/net/pktgen/pgctrl
```

	rx	tx
bytes	137.87 MiB	0 KiB
max	334.60 Mbit/s	0 kbit/s
average	59.45 Mbit/s	0.00 kbit/s
min	0 kbit/s	0 kbit/s
packets	2409506	0
max	713824 p/s	0 p/s
average	126816 p/s	0 p/s
min	0 p/s	0 p/s

Need reliable ways and tunings
to measure performance

Notes on this experiment

- No guarantee as usual
- No tuning effort has been made
- References:
 - <http://www.intel.com/go/dpdk>
 - For pktgen (in Japanese)
 - <http://research.sakura.ad.jp/2010/10/08/infini01/>